Big Data Video Presentation

Question d:

Uploading CSV in S3

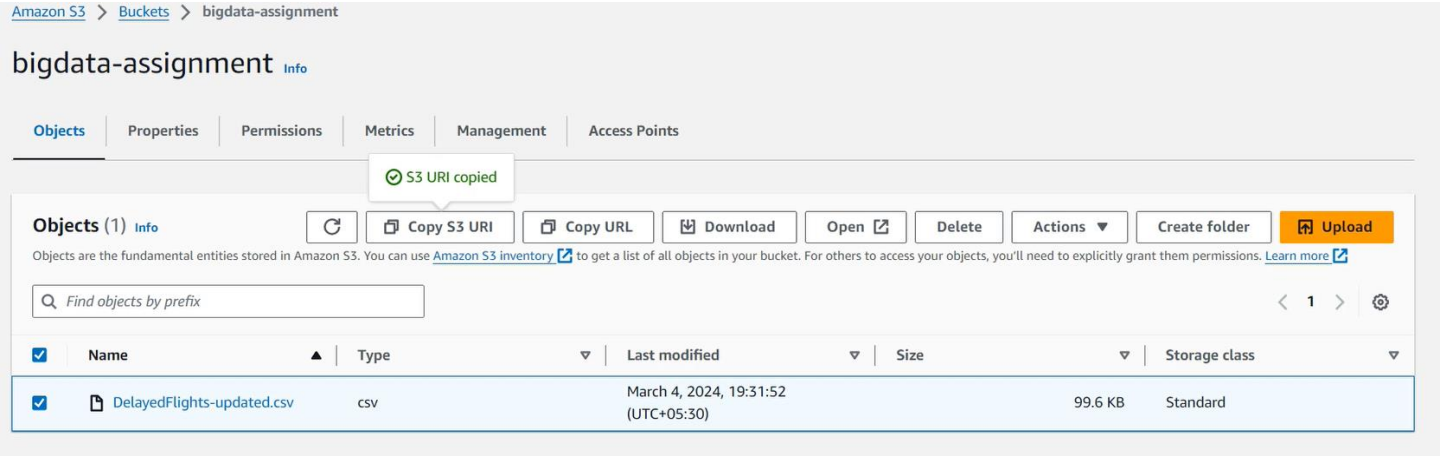

Spark Initialization

```
Welcome to
      ____              __
     / __/__  ___ _____/ /__
    _\ \/ _ \/ _ `/ __/  '_/
   /__ / .__/\_,_/_/ /_/\_\   version 3.5.0-amzn-0
      /_/

Using Python version 3.9.16 (main, Sep  8 2023 00:00:00)
Spark context Web UI available at http://ip-172-31-55-229.ec2.internal:4040
Spark context available as 'sc' (master = yarn, app id = application_1709559670189_0001).
SparkSession available as 'spark'.
>>> from pyspark.sql import SparkSession
>>> spark = SparkSession.builder.appName("delay_analysis").getOrCreate()
24/03/04 13:59:17 WARN SparkSession: Using an existing Spark session; only runtime SQL configurations will take effect.
>>> file_path ='s3://bigdata-assignment/DelayedFlights-updated.csv'
>>> df = spark.read.option("header", "true").csv(file_path)
>>> df.createOrReplaceTempView("delay_flights")
24/03/04 14:03:22 WARN SparkStringUtils: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.
>>> result = spark.sql("SELECT Year, AVG((CarrierDelay / ArrDelay) * 100) AS AverageDelayPercentage FROM delay_flights GROUP BY Year")
>>> result.show()
+----+----------------------+
|Year|AverageDelayPercentage|
+----+----------------------+
|2010|     21.89310246015957|
|2003|    24.557549755575373|
|2005|     28.01977637202288|
|2009|     28.33058554239575|
|2006|    30.453296261292596|
|2004|     43.64459443230066|
|2008|     28.88346981456985|
|2007|    19.850007017971283|
+----+----------------------+
```

Spark Code

```
from pyspark.sql import SparkSession

# Create a Spark session

spark = SparkSession.builder.appName("DelayFlightsAnalysis").getOrCreate()


# Upload CSV and create DataFrame

csv_path = "path/to/DelayedFlights-updated.csv"

delay_flights_df = spark.read.csv(csv_path, header=True, inferSchema=True)


# Create a temporary table named delay_flights

delay_flights_df.createOrReplaceTempView("delay_flights")
```

```
# Run the first query using HiveQL
result_hiveql = spark.sql("""
   SELECT Year, AVG((CarrierDelay / ArrDelay) * 100) AS avg_delay_percentage
   FROM delay_flights
   GROUP BY Year""").show()

# Run the second query using Spark SQL
result_sparksql = spark.sql("""
   SELECT Year, AVG((CarrierDelay / ArrDelay) * 100) AS avg_delay_percentage
   FROM delay_flights
   GROUP BY Year""").show()
```

Question e:

EMR Cluster Initialization

```
C:\Users\tuanm\OneDrive\Desktop\DS\bigdata\new>ssh -i Bigdata.pem hadoop@ec2-52-91-134-150.compute-1.amazonaws.com
   ,     #_
   ~\_  ####_        Amazon Linux 2023
  ~~  \_#####\
  ~~     \###|
  ~~       \#/ ___   https://aws.amazon.com/linux/amazon-linux-2023
   ~~       V~' '->
    ~~~         /
      ~~._.   _/
         _/ _/
       _/m/'
Last login: Mon Mar  4 14:04:25 2024

EEEEEEEEEEEEEEEEEEE MMMMMMMM           MMMMMMMM RRRRRRRRRRRRRRR
E::::::::::::::::::E M:::::::M         M:::::::M R::::::::::::::R
EE:::::EEEEEEEEE:::E M::::::::M       M::::::::M R:::::RRRRRR:::::R
  E::::E       EEEEE M:::::::::M     M:::::::::M RR::::R      R::::R
  E::::E             M::::::M:::M   M:::M::::::M   R:::R      R::::R
  E:::::EEEEEEEEEE   M:::::M M:::M M:::M M:::::M   R:::RRRRRR:::::R
  E::::::::::::::E   M:::::M  M:::M:::M  M:::::M   R:::::::::::RR
  E:::::EEEEEEEEEE   M:::::M   M:::::M   M:::::M   R:::RRRRRR:::R
  E::::E             M:::::M    M:::M    M:::::M   R:::R      R::::R
  E::::E       EEEEE M:::::M     MMM     M:::::M   R:::R      R::::R
EE:::::EEEEEEEE:::E M:::::M             M:::::M   R:::R      R::::R
E::::::::::::::::::E M:::::M             M:::::M RR::::R      R::::R
EEEEEEEEEEEEEEEEEEE MMMMMMM             MMMMMMM RRRRRRR      RRRRRR
```

Adding Shell Script and table in Hadoop directory

```
C:\Users\tuanm\OneDrive\Desktop\DS\bigdata\new>scp -i Bigdata.pem DelayedFlights-updated.csv hadoop@ec2-52-91-134-150.compute-1.amazonaws.com:/home/hadoop/
DelayedFlights-updated.csv                                    100%  100KB   1.6KB/s   01:04

C:\Users\tuanm\OneDrive\Desktop\DS\bigdata\new>scp -i Bigdata.pem assignment.sh hadoop@ec2-52-91-134-150.compute-1.amazonaws.com:/home/hadoop/
assignment.sh                                                 100% 3557    11.0KB/s   00:00
```

Grant permission to Shell Script and execute

```
[hadoop@ip-172-31-55-229 ~]$ chmod +x assignment.sh
[hadoop@ip-172-31-55-229 ~]$ ./assignment.sh
Hive Session ID = 6a8f9b0d-62d1-4063-b0fe-e8df4f23f963
```

Map Reducer Worker distribution

```
Map 1: -/-      Reducer 2: 0/1
Map 1: 0/1      Reducer 2: 0/1
Map 1: 0/1      Reducer 2: 0/1
Map 1: 0(+1)/1  Reducer 2: 0/1
Map 1: 0/1      Reducer 2: 0/1
Map 1: 1/1      Reducer 2: 0(+1)/1
Map 1: 1/1      Reducer 2: 1/1
OK
```

Run Time Calculation

```
Time taken: 23.629 seconds, Fetched: 1 row(s)
Iteration 1
HiveQL Execution Time: 0m43.761s seconds
Spark SQL Execution Time: 0m44.568s seconds
Spark SQL Execution Time: 0m44.568s seconds
Iteration 2
HiveQL Execution Time: 0m45.451s seconds
Spark SQL Execution Time: 0m44.386s seconds
Iteration 3
HiveQL Execution Time: 0m41.443s seconds
Spark SQL Execution Time: 0m44.888s seconds
Iteration 4
HiveQL Execution Time: 0m44.945s seconds
Spark SQL Execution Time: 0m44.805s seconds
Iteration 5
HiveQL Execution Time: 0m39.552s seconds
Spark SQL Execution Time: 0m49.065s seconds
```

Download into local path

```
C:\Users\tuanm\OneDrive\Desktop\DS\bigdata\new>scp -i Bigdata.pem hadoop@ec2-52-91-134-150.compute-1.amazonaws.com:execution_times.csv .
execution_times.csv
```

Results

	Iteration	HiveQL_Time	SparkSQL_Time
0	1	47.165	50.591
1	2	44.515	46.060
2	3	40.606	43.760
3	4	43.132	41.166
4	5	43.893	44.501

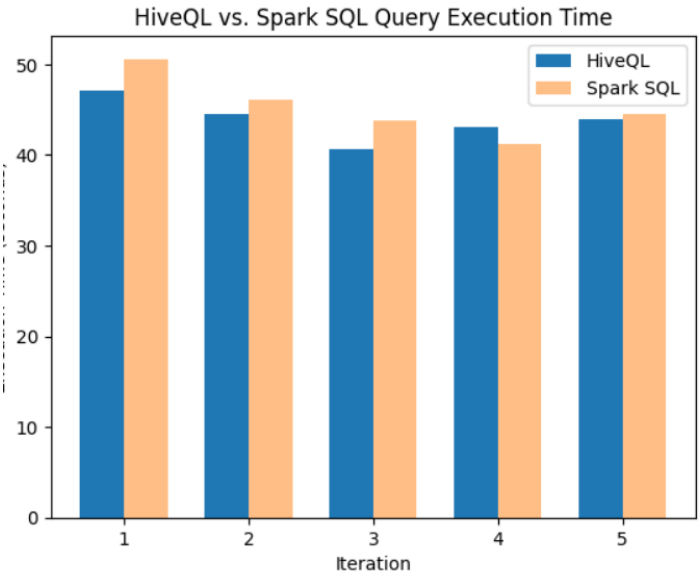

Shell Script file is **assignment_for_5_iteration.sh**

Question f:

Same command approach as above. But Shell script is different. I have attached it in the same folder
assignment_for_5_queries.sh

Results

	Mapping	HiveQL_Time	SparkSQL_Time
0	Career delay query	58.340	43.824
1	Nas delay query	42.040	41.719
2	Weather delay query	43.682	44.415
3	Late aircraft delay query	42.455	43.383
4	Security delay query	34.011	44.908

Question g:

Results